

Supervised Algorithms

Regression

- Single and multivariate
- $h_{\mathbf{w}}(\mathbf{x}_j) = \mathbf{w}^T \mathbf{x}_j = \sum_i w_i x_{j,i}$
- $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \sum_j L_2(y_j, \mathbf{w}^T \mathbf{x}_j)$
- Regularization γ to penalize complexity
- $Cost(h) = EmpLoss(h) + \gamma L_q(\mathbf{w})$
- Complexity: $L_q(\mathbf{w}) = \sum_i |w_i|^q$
- **Linear Regression**
 - Directly computable
- **Non-Linear Regression**
 - Differentiable, use *gradient descent*
- **Stepwise Regression** (Perceptron function)
 - Linearly separable data/Hard Threshold
 - $h_{\mathbf{w}}(\mathbf{x}) = Threshold(\mathbf{w} \cdot \mathbf{x})$
 - $Threshold(z) = 1$ if $z \geq 0, 0$ otherwise
 - Not differentiable - use perceptron learning
- **Logistic Regression** (Sigmoid Function)
 - Soft threshold
 - $Logistic(z) = \frac{1}{1+e^{-z}}$
 - $h_{\mathbf{w}}(\mathbf{x}) = g(\mathbf{w} \cdot \mathbf{x}) = Logistic(\mathbf{w} \cdot \mathbf{x})$
 - Differentiable, use *gradient descent*
 - Nice property: $g'(u) = g(u)(1 - g(u))$

Decision Trees (DT)

- *Analogy*: human decision-making paradigm
- Can represent any truth table
- *Inductive bias*: best attribute/dimension first
- Best = Highest information gain (entropy)
- Hypothesis complexity: # tree levels and nodes
- Discrete attributes: split point for each value
- Continuous attributes: establish split threshold
- Prune tree to eliminate overfitting
- **Iterative Dichotomiser (ID3)**

Instance Based

- Nonparametric: grows with # of examples
- *Inductive Bias*:
 - locality (near points are similar)
 - all dimensions are of equal importance
- Lazy learner/just-in-time learning
- Curse of Dimensionality, w/ N examples, d dim.:
 - $N = O(2^d)$ to maintain generalization accuracy.
- **K Nearest Neighbors (KNN) Classification**
 - Given x_q , find k nearest examples
 - Classify based on plurality vote w/ odd k
 - Low k can overfit, and high k underfit
 - Minkowski distance:
$$L_p(x_q, x_j) = (\sum_i |x_{j,i} - x_{q,i}|^p)^{1/p}$$
 - $p = 2$: Euclidean Distance
 - $p = 1$: Manhattan ""
 - Any other abstract distance function
 - Normalize if dimension scale changes
- **KNN average regression**
 - $h(x) = \sum_i y_i / k$, $y_i \in k$ nearest points.

- larger k : smoother spikes; discontinuities
- **KNN linear regression**
 - Plot a line through k examples
 - Discontinuities still an issue
- **Locally weighted regression**
 - Smooth curve, avoid discontinuities
 - Use kernel function $K(dist(x_q, x_j))$
 - more weight on closer points
 - $w^* = \operatorname{argmin}_w \sum_j K(d(x_q, x_j))(y_j - w \cdot x_j)^2$
 - $h(x_q) = w^* \cdot x_q$
 - Must solve for w^* for \forall query points.
 - But only points overlapping kernel matter.

Ensemble Learning

- Vote among multiple hypotheses
- *Inductive Bias*: Mult. hyp. generalize better
- $\Pr\{h_k(x) \neq y_k\}$ indep for $\forall k$
- more complex h via linear combo of h_k .
- **Boosting/AdaBoost**
 - For each h_i (weak classifier)
 - correctly classified x : decrease weight
 - incorrectly classified: increase weight
 - Ultimately more emphasis on misclassified
 - Produce weighted average of all hypotheses
 - $f(x) = \sum_{t=1}^k \alpha_t h_t(x)$, a strong classifier
 - $\uparrow k \rightarrow$ near perfect training error
 - requires $h(i)$ has error rate $> 50\% + \epsilon$
 - Will overfit if weak learner always overfits
- **Random Forest**
 - *Analogy*: multiple interviewers
 - Randomly ask different questions
 - Vote on whether to hire candidate
 - Multiple decision trees vote on answer
 - Random sampling prevents overfitting
 - For each of k trees:
 - Randomly sample from a subset of X
 - "" "" subset of features
 - Learn tree from the sampled data

Neural Networks (NN)

- *Analogy*: brain neural connections and learning
- Input layer, $0 \leq$ hidden layers, output layer
- Domain knowledge helps establish structure
- Model *any* function with enough nodes/layers
- For each layer i w/ n_i inputs $\mathbf{a}_i \in \mathbb{R}^{n_i}$
 - Bias weight $b_i \in \mathbb{R}$
 - Input weights to next layer: $W_i \in \mathbb{R}^{n_i \times n_{i+1}}$
 - Activation function g
 - Apply to weighted sum of inputs + bias
 - $\mathbf{a}_{i+1} = g(W_i^T \mathbf{a}_i + b_i)$
- Perceptron (step) activation function
 - $g(z) = 1$ if $z > 0, 0$ otherwise
- Logistic/Sigmoid Function (Gaussian integrated):
 - $g(z) = \frac{1}{1+e^{-z}}$
- **Back Propagation** learning
 - Initiate input layer weights and bias
 - Observed error: $L_2(W) = |\mathbf{y} - h_W(\mathbf{x})|^2$
 - Update output weights via observed error

- Use *gradient descent*
- From output to earliest hidden layer:
 - Propagate Δ values to previous layer
 - Update weights between the two layers
- (each hidden node resp. for some error)

Deep Learning

- Applications:
 - B/W image colorization
 - Sound generation from silent media
 - Machine translations
 - Object detection
 - Handwriting generation and recognition
 - [Eloquent] Text generation via recurrent NN
 - Game Playing (in a human-like way)
 - Speech recognition
 - Image object detection, classification, de-abstractization, sketch inversion
 - Natural Language Processing

Support Vector Machines (SVM)

- Maximum margin separator between examples
- Non-param., but needs frac. of ex. (support vectors)
- Expensive: use quadratic programming - $O(n^3)$
- Support vectors lie along decision boundary
- **Kernels** (kernel function)
 - notion of similarity between data
 - linearly-separable data:
 - linear kernel: $K(x, y) = x^T y$
 - non-linearly separable data
 - polynomial: $K(x, y) = (x^T y + c)^d$
 - project data to higher dimension
 - dividing hyperplane in $d + 1$ dimensions
 - d degrees of freedom
 - Radial-Basis: $K(x, y) = e^{-(\|x-y\|^2/(2\sigma^2))}$
 - Sigmoid: $K(x, y) = \tanh(\alpha x^T y + \theta)$
 - Kernels for other abstract similarity...
 - Requires some domain knowledge of data
- Regularization constant C and Soft-Margin SVM
 - $\downarrow C$: permit points to fall inside margin
- Multi-class classification: 1vs1 or 1vs \forall strategies

Randomized Algorithms/Optimization

- Given i/p X , objective/fitness func. $f: X \mapsto \mathbb{R}$
 - find $x \in X$ such that $f(x) = \max_x f(x)$
- **Hill Climbing (HC)**
 - *Analogy*: Climb Everest in fog w/ amnesia
 - Start at arbitrary location in input space
 - Move to the highest-value neighbor
 - If neighbor $>$ current, proceed to neighbor
 - Else return current. Possibly only *local max*.
- **Random Restart Hill Climbing (RHC)**
 - Alleviates local max somewhat
 - Restarts at random point a constant # of times.
- **Simulated Annealing (SA, Metropolis Hastings)**
 - *Analogy*: Rept'd heat/cool'g strengthens blade.
 - Allow bad moves, but with decreasing freq.

- T : some gradually decreasing temperature func.
- $\Delta E = \text{Value}(\text{next}) - \text{Value}(\text{curr})$
- If $\Delta E \leq 0$ take bad move w/ prob. $e^{\Delta E/T}$
- As $T \rightarrow 0$, transitions *Random Walk* \rightarrow *HC*.

- Genetic Algorithms (GA)

- *Analogy*: Natural selection and mutation
- *Apps*: Opt problems w/ approp. encoding
 - 8-Queens, circuit layout, job schedule
- Select most fit pairs among population
- Reproduce (cross-over) each pair
 - One-point crossover strategy
 - Uniform crossover strategy (random bits)
- Mutate offspring with small probability
- Replace least-fit individuals with new offspring
- Repeat until convergence

- MIMIC

- Model probability distribution instead of:
- Population (non-parametric representation)
- Convey structure
- Generate samples from distribution $P^{\theta_t}(x)$
- Set θ_{t+1} to n 'th percentile
- Retain only samples with $f(x) \geq \theta_{t+1}$
- Estimate $P^{\theta_{t+1}}(x)$
- Repeat until convergence
- Technical details:
 - Estimate distribution via dependency tree
 - Use the *KL divergence* from info theory
 - Vastly fewer iterations than above algs
 - Each iteration much more costly

Bayesian

- Most prob. h given the data: $\arg\max_{h \in H} \Pr\{h|D\}$
- Bayes Rule: $\Pr\{h|D\} = \Pr\{D|h\} \Pr\{h\} / \Pr\{D\}$
- Maximum a posteriori (MAP):
 - $h_{MAP} = \arg\max_h \Pr\{D|h\} \Pr\{h\}$
 - disregard the normalizer $\Pr\{D\}$
- Maximum likelihood: $h_{ML} = \arg\max_h \Pr\{D|h\}$
 - assumes uniform $\Pr\{h\}$. Actual h irrelevant.
 - if Gaussian noise in data $\rightarrow \arg\min_h L_2$ error.
- Problems:
 - Requires domain knowledge of prior probabilities
 - Expensive to compute, being linear in $|H|$
- **Bayesian Optimal Classifier**
 - Most probable *classification* given the data
 - Combine weighted predictions of $\forall h_i$ given data
 - $V_{MAP} = \arg\max_v \sum_h \Pr\{v|h\} \Pr\{h|D\}$
- **Bayesian Belief Networks** (Bayes Nets)
 - Cond. indep. assumptions for X and Y given Z
 - $\Pr\{X|Y, Z\} = \Pr\{X|Z\}$
 - Conditional probabilities
 - Sample nodes in *Topological Sort*
 - *apps*: simulate/approximate a complex process
- **Naive Bayes**
 - assume data independence given parent.
 - \rightarrow assume attr. independence given class.
 - No guarantee in the real world.
 - *apps*: classifying text documents
 - Attributes: words. Values: frequencies, OR
 - Attributes: word positions. Values: words.

- $V_{NB} = \arg\max_V \Pr\{V\} \prod_i \Pr\{a_i|V\}$
- Requires sufficiently large set of training data
- Tractable: requires no search

Classification Metrics

- Accuracy: $A(y, y^*) = (1/N) \sum_{i=0}^N 1(y_i^* = y_i)$
- Binary Classifications (t_p = true positive)
 - Precision: $P(y, y^*) = |y \cap y^*|/|y| = t_p/(t_p + f_p)$
 - Recall: $R(y, y^*) = |y \cap y^*|/|y^*| = t_p/(t_p + f_n)$
 - $F_\beta(y, y^*) = (1 + \beta^2) \frac{P(y, y^*)R(y, y^*)}{\beta^2(P(y, y^*) + R(y, y^*))}$
- Multiclass Classifications
 - $M_B(y, y^*)$: any binary scoring metric, L : labels
 - Macro: apply equal weight to each class
 - $(1/|L|) \sum_{l \in L} M(y_l, y_l^*)$
 - Weighted: average based on class presence
 - $\frac{1}{\sum_{l \in L} |y_l^*|} \sum_{l \in L} |y_l^*| M(y_l, y_l^*)$
 - Micro
 - None: $\{M(y_l, y_l^*) | l \in L\}$ (metric for \forall labels)

Unsupervised Algorithms

Clustering

- **Single-Linkage Clustering**
 - Start with n points
 - Intercluster distance: closest two points in each
 - Alternatives: median, mean distance
 - Merge two closest clusters
 - Repeat $n - k$ times to produce k clusters
 - $O(n^3)$, but practically fast
- **K-Means Clustering** (in Euclidean Space)
 - 'Hard' clustering (special case of EM Clustering)
 - *Apps*: Image segmentation and compression
 - Pick k centers at random (or distributed)
 - Each center "claims" its closest points
 - Closest \rightarrow minimizing L_2 error.
 - Recompute the centers (avg clustered points)
 - Repeat until convergence (to *local minimum*!)
 - $O(kn)$ per iteration, and $O(k^n)$ iterations.
- **K-Medoids Clustering**
 - Not Euclidean-Space \rightarrow cannot use L_2 error.
 - Use abstract $\mathcal{V}(x_n, \mu_k)$ instead of $\|x_n - \mu_k\|^2$
 - Can't take average of clusters.
 - Can assign μ_k to one of cluster points
 - $O(kn + n_k^2)$ per iteration
- **Expectation-Maximization** (EM Clustering)
 - 'Soft' clustering
 - Point shared by mult. clusters prob'ly
 - Gaussian mixture model:
 - $f(x) = \sum_{i=1}^k \pi_i N_i(x|\mu_i, \sigma_i^2)$
 - π_i : mixing coefficient. $\sum_{i=1}^k \pi_i = 1$
 - Mixed Gaussian (μ, σ^2) for each cluster
 - Pick initial hidden vars: μ_k , covariances, and π_k .
 - E - *step*: prob. of component k explaining x
 - M - *step*: Re-estimate hidden vars
 - Iterate until log-likelihood reaches convergence
 - $\log \Pr\{\mathbf{X}|\mu, \sigma^2, \pi\} =$

- $\sum_{n=1}^N \log \left\{ \sum_{k=1}^K \pi_k N(x_n|\mu_k, \sigma_k^2) \right\}$
- Local optima possible \rightarrow random restart.

Dimensionality Reduction

- *Filtering*: Feature selection abstracted from learner
 - Fast, but ignore the learning problem
 - Ex: DT learning, information gain, variance
- *Wrapping*: Coupled with learner as one routine
 - Takes model bias into account, but SLOW.
 - Ex: Hill Climbing, Rand. Algorithms
- *Forward Search*:
 - Start with empty set of features.
 - Add most 'contributing' features until threshold.
- *Backward Search*:
 - Start with full set of features.
 - Remove most 'useless' features until threshold.
- **Principle Component Analysis** (PCA)
 - Mutually orthogonal and ordered transformed features
 - Gravitate towards 'average' of features
 - 'Best' dimension: highest *variance/eigenvalue*
 - This is the *First Principle Component*
 - Find via *SVD* in Linear Algebra
 - Find orthogonal (*Second Principle*) component
 - Idea: min. L_2 error moving from $N \rightarrow M$ dims
 - Select best components with non increasing var.
 - Eliminate dimensions beyond best M
- **Independent Component Analysis** (ICA)
 - Mutually independent transformed features
 - Gravitate towards 'abstracted' components
 - *Apps*: Cocktail party problem, mixed models
 - Min. *mutual information* for transf. features
 - Maximizing *kurtosis* of a dim. is one way.
 - $\rightarrow I(y_i; y_j) = 0$ (or minimum)
 - $I(X; Y)$ is max. (between *orig.* & transformed)
 - Idea: given observables, find indep. hidden var.
- **Random Component Analysis** (RCA)
 - Linear transformation in random directions
 - *Random* linear. combo of orig. dims still useful
 - For $N \rightarrow M$, M normally \uparrow than PCA/ICA.
 - Fast. But requires many random trials.
- **Linear Discriminant Analysis** (LDA)
 - Finds projection that *discriminates* on label.

Classification Metrics

Reinforcement Learning

Markov Decision Processes (MDP)

- *Markov*: Only present matters. Stationary rules.
- $a \in \text{Actions}$, $s, s' \in \text{States}$
- *Model*: $T(s, a, s') \approx \Pr\{s'|s, a\}$
- *Reward*: $R(s), R(s, a), R(s, a, s')$ (typically 1st form)
- *Policy*: $\pi(s) \rightarrow a$.
 - π^* : optimal policy w/ max reward.
- *Infinite horizon*: navigate world forever, ∞ rewards.
- *Finite horizon*: discount γ to incentivise finish.

- $\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') U(s')$
- **Bellman Equation:**
 - $U(s) = R(s) + \gamma \max_a \sum_{s'} T(s, a, s') U(s')$
 - $U(s)$: state utility, $R(s)$: state reward
 - n states $\rightarrow n$ equations, n , unknowns
 - non-linear due to \max operation
- **Value Iteration (VI)**
 - Start w/ arbitrary utilities
 - Update based on neighbors
 - $\hat{U}_{t+1}(s) = R(s) + \gamma \max_a \sum_{s'} T(s, a, s') \hat{U}_t(s')$
 - Repeat until convergence
 - Then solve for π^* , straightforward
 - Lends itself for parallel computation
- **Policy Iteration (PI)**
 - Start with arbitrary π_0
 - Given π_t , calculate $U_t = U^{\pi_t}$ (follow policy)
 - $U_t(s) = R(s) + \gamma \sum_{s'} T(s, \pi_t(s), s') U_t(s')$
 - The action is fixed from the policy $\pi_t(s)$
 - n linear eq, n unknowns. \rightarrow Linear Algebra.
 - Improve: $\pi_{t+1} = \operatorname{argmax}_a \sum_{s'} T(s, a, s') U_t(s')$

Reinforcement Learning (model-free)

- No model (transition probabilities) or rewards
- Given transitions $\langle s, a, r, s' \rangle$, learn policy
- **Q-Learning**
 - $Q(s, a) = R(s) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q(s', a')$
 - Derive U and π from Q :
 - $U(s) = \max_a Q(s, a)$
 - $\pi(s) = \operatorname{argmax}_a Q(s, a)$
 - Estimate from transitions (w/out R & T):
 - $Q(s, a) \xleftarrow{\alpha_t} r + \gamma \max_{a'} Q(s', a')$
 - α_t : learning rate
 - $\max_{a'} Q(s', a')$: util. of next state
 - Notation: $v \xleftarrow{\alpha} x = v \xleftarrow{(1-\alpha)} v + \alpha x$
 - \hat{Q} starts anywhere (results will vary)
 - Following update rule above, $\hat{Q}(s, a) \rightarrow Q(s, a)$
 - Visiting $s, a \infty$ times, $s' \approx T(s, a, s'), r \approx R(s)$
 - Questions:
 - Initial \hat{Q}
 - α_t decay factor?
 - action choosing policy?
 - Take random action *sometimes*
 - Otherwise take best action
 - Mimics simulated annealing
 - *Exploration vs. Exploitation*
 - Exploration: cont. to learn Q
 - Exploitation: quicker maximize π^*

Game Theory

Zero-Sum Games

- **Minimax/Maximin**
 - Player A considers worse-case strategy by B
 - A chooses maximum minimum value by B
 - B chooses the minimum maximum value by A
 - Applies to *perfect* or *hidden* information games

- *Von Neumann*
 - 0-sum games of per. info: *minimax*=*maximin*.
 - \exists optimal strategy for each player
- Optimal *Mixed Strategy* for 2X2 game
 - Two players, A & B, strategies A_1, A_2, B_1, B_2
 - Mixed strategy gains: $\mathbf{m}_{AB} : m_{11}, m_{21}, m_{12}, m_{22}$
 - A uses strategy 1 w/ probability p
 - eg_i : Expected gain for A given B uses strategy i
 - $eg_1 = m_{11}p + m_{21}(1-p)$
 - $eg_2 = m_{12}p + m_{22}(1-p)$
 - $p^* = \max_p \{ \min \{ eg_1, eg_2 \} \}$
 - Either $p = 0, p = 1$, or where both equal
- Optimal *Mixed Strategy* for $n \times m$ game
 - $\mathbf{P}_A = \{p_1, p_2, \dots, p_n\}$
 - $eg_j = \sum_{i=1}^n m_{ij} p_i$ for all $j \in 1, \dots, m$
 - n unknowns, m eq. \rightarrow Linear Programming
 - Choose \mathbf{P}_A to maximize $\min \{ eg_1, \dots, eg_m \}$
 - such that $\sum p_i = 1, 0 \leq p_i \leq 1, \forall i$

Non-Zero-Sum Games

- Non-deterministic, non-cooperative, hidden info
- Typical example: Prisoner's Dilemma
- n strategy spaces $\mathbf{S}_1, \mathbf{S}_2, \dots, \mathbf{S}_n$
- n payoff functions u_1, \dots, u_n st $u_i : \mathbf{S}_1 \times \dots \times \mathbf{S}_n \rightarrow \mathbb{R}$
- What's the optimal mixed strategy?
- **Nash Equilibrium (NE)**
 - $s_1^*, \dots, s_n^* \in S_1 \times \dots \times S_n$ are a NE iff
 - $\forall i. s_i^* = \operatorname{argmax}_{s_i} u_i(s_1^*, \dots, s_i, \dots, s_n^*)$
 - ($\neg \exists$ reason for any one player to switch)
 - Works for pure & mixed strategies
 - Assumes best action for self (regardless of others)
 - Can contain multiple Nash Equilibria
 - An *implausible* threat hinders own utility
 - n repeated games $\rightarrow n$ repeated NE
 - assuming implausible threats

Bayesian Games

- Action spaces: $\mathbf{A}_1, \dots, \mathbf{A}_n$
- Type spaces: $\mathbf{T}_1, \dots, \mathbf{T}_n$
- Beliefs: $\mathbf{P}_1, \dots, \mathbf{P}_n$
- $P_{-i}(t_{-i}|t_i) = \text{PDF of others' types given own}$
- Payoff functions: u_1, \dots, u_n
- $u_i(a_1, \dots, a_n, t_i)$ - payout to player i w/ type i
 - player j chooses action a_j for $\forall j$
- All players know their own $\mathbf{A}_i, \mathbf{T}_i, \mathbf{P}_i, u_i$
- Strategy $S_i : \mathbf{T}_i \rightarrow \mathbf{A}_i$
- **Bayesian Nash Equilibrium (BNE)**
 - $s_1^*, \dots, s_n^* \in S_1 \times \dots \times S_n$ are a BNE iff $\forall i$ and $\forall t_i \in \mathbf{T}_i. s_i^*(t_i) = \operatorname{argmax}_{a_i \in \mathbf{A}_i} \{ \sum_{t_{-i} \in \mathbf{T}_{-i}} u_i(s_1^*(t_1), \dots, a_i, \dots, s_n^*(t_n)) \times P_i(t_{-i}|t_i) \}$

Repeated games w/ uncertain end

- # rounds left is uncertain

- $\Pr \{ \text{play again} \} = \gamma, \Pr \{ \text{game over} \} = 1 - \gamma$
- $E[\# \text{ rounds}] = 1/(1 - \gamma)$
- **Tit for Tat** strategy
 - Cooperate 1st round, copy opponent move after
- **Grim Trigger** strategy
 - Cooperate while opponent cooperates.
 - Once line is crossed, forever defect
- **Pavlov** strategy
 - Cooperate if agree, defect if disagree.
 -
- *Subgame perfect (SP)*
 - Str. is always best response indep. of history
 - $\neg \text{SP}$ if \exists history of moves st. str. implausible
 - $\neg \text{SP} \iff \exists$ pair of str's not leading back to mutual cooperate
- *Mini-Max Profile* (For zero-sum game)
 - Min guaranteed payoffs for \forall players on defense
 - Applicable in pure or mixed strategies

Stochastic Games & Multiagent RL

- *Analogy*: MDP-RL::Stochastic game-Multiagent RL
- More general than MDPs or other previous models
- **S**: states
- **A_i**: actions for player $i, a, b, a \in \mathbf{A}_1, b \in \mathbf{A}_2$
- **T**: transition probabilities $T(s, (a, b), s')$
- **R_i**: rewards for payer $i, R_1(s, (a, b)), R_2(s, (a, b))$
- γ : discount
- Impose restrictions to produce other models:
 - $R_1 = -R_2 \rightarrow$ 0-sum stochastic game
 - $T(s, (a, b), s') = T(s, (a, b'), s'), R_2(s, (a, b)) = 0, R_1(s, (a, b)) = R_1(s, (a, b'))$
 - $\forall b' \rightarrow$ MDP, makes second player irrelevant
 - $|\mathbf{S}| = 1 \rightarrow$ repeated games
- **Zero-Sum Stochastic Games**
 - $Q_i^*(s, (a, b)) = R_i(s, (a, b)) + \gamma \sum_{s'} T(s, (a, b), s') \min_{a', b'} Q_i^*(s', (a', b'))$
 - *Q-Learning* for transition $\langle s, (a, b), (r_1, r_2), s' \rangle$
 - $Q_i(s, (a, b)) \xleftarrow{\alpha} r_i + \gamma \min_{a', b'} Q_i(s', (a', b'))$
 - Value Iteration works
 - Minimax-Q converges
 - Unique solution to Q^*
 - Policies can be computed independently
 - Update efficient (polytime)
 - Q functions sufficient to specify policy
- **General-Sum Stochastic Games (Nash-Q)**
 - Same as Q_i^*/Q -Update above, but use Nash Equilibrium instead of minimax/maximin
 - Value Iteration doesn't work
 - Nash-Q doesn't work
 - No unique solution
 - Policies can't be computed independently
 - Update not efficient
 - Q functions not sufficient to specify policy