

# Supervised Algorithms

## Regression

- Single and multivariate
- $h_{\mathbf{w}}(\mathbf{x}_j) = \mathbf{w}^T \mathbf{x}_j = \sum_i w_i x_{j,i}$
- $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \sum_j L_2(y_j, \mathbf{w}^T \mathbf{x}_j)$
- Regularization  $\gamma$  to penalize complexity
- $Cost(h) = EmpLoss(h) + \gamma L_q(\mathbf{w})$
- Complexity:  $L_q(\mathbf{w}) = \sum_i |w_i|^q$
- **Linear Regression**
  - Directly computable
- **Non-Linear Regression**
  - Differentiable, use *gradient descent*
- **Stepwise Regression** (Perceptron function)
  - Linearly separable data/Hard Threshold
  - $h_{\mathbf{w}}(\mathbf{x}) = Threshold(\mathbf{w} \cdot \mathbf{x})$
  - $Threshold(z) = 1$  if  $z \geq 0, 0$  otherwise
  - Not differentiable - use perceptron learning
- **Logistic Regression** (Sigmoid Function)
  - Soft threshold
  - $Logistic(z) = \frac{1}{1+e^{-z}}$
  - $h_{\mathbf{w}}(\mathbf{x}) = g(\mathbf{w} \cdot \mathbf{x}) = Logistic(\mathbf{w} \cdot \mathbf{x})$
  - Differentiable, use *gradient descent*
  - Nice property:  $g'(u) = g(u)(1 - g(u))$

## Decision Trees (DT)

- *Analogy*: human decision-making paradigm
- Can represent any truth table
- *Inductive bias*: best attribute/dimension first
- Best = Highest information gain (entropy)
- Hypothesis complexity: # tree levels and nodes
- Discrete attributes: split point for each value
- Continuous attributes: establish split threshold
- Prune tree to eliminate overfitting
- **Iterative Dichotomiser (ID3)**

## Instance Based

- Nonparametric: grows with # of examples
- *Inductive Bias*:
  - locality (near points are similar)
  - all dimensions are of equal importance
- Lazy learner/just-in-time learning
- Curse of Dimensionality, w/  $N$  examples,  $d$  dim.:
  - $N = O(2^d)$  to maintain generalization accuracy.
- **K Nearest Neighbors (KNN) Classification**
  - Given  $x_q$ , find  $k$  nearest examples
  - Classify based on plurality vote w/ odd  $k$
  - Low  $k$  can overfit, and high  $k$  underfit
  - Minkowski distance:
$$L_p(x_q, x_j) = (\sum_i |x_{j,i} - x_{q,i}|^p)^{1/p}$$
    - $p = 2$ : Euclidean Distance
    - $p = 1$ : Manhattan ""
    - Any other abstract distance function
    - Normalize if dimension scale changes
- **KNN average regression**
  - $h(x) = \sum_i y_i / k$ ,  $y_i \in k$  nearest points.

- larger  $k$ : smoother spikes; discontinuities
- **KNN linear regression**
  - Plot a line through  $k$  examples
  - Discontinuities still an issue
- **Locally weighted regression**
  - Smooth curve, avoid discontinuities
  - Use kernel function  $K(dist(x_q, x_j))$ 
    - more weight on closer points
  - $w^* = \operatorname{argmin}_w \sum_j K(d(x_q, x_j))(y_j - w \cdot x_j)^2$
  - $h(x_q) = w^* \cdot x_q$
  - Must solve for  $w^*$  for  $\forall$  query points.
  - But only points overlapping kernel matter.

## Ensemble Learning

- Vote among multiple hypotheses
- *Inductive Bias*: Mult. hyp. generalize better
- $\Pr\{h_k(x) \neq y_k\}$  indep for  $\forall k$
- more complex  $h$  via linear combo of  $h_k$ .
- **Boosting/AdaBoost**
  - For each  $h_i$  (weak classifier)
    - correctly classified  $x$ : decrease weight
    - incorrectly classified: increase weight
  - Ultimately more emphasis on misclassified
  - Produce weighted average of all hypotheses
  - $f(x) = \sum_{t=1}^k \alpha_t h_t(x)$ , a strong classifier
  - $\uparrow k \rightarrow$  near perfect training error
  - requires  $h(i)$  has error rate  $> 50\% + \epsilon$
  - Will overfit if weak learner always overfits
- **Random Forest**
  - *Analogy*: multiple interviewers
    - Randomly ask different questions
    - Vote on whether to hire candidate
  - Multiple decision trees vote on answer
  - Random sampling prevents overfitting
  - For each of  $k$  trees:
    - Randomly sample from a subset of  $X$
    - "" "" subset of features
    - Learn tree from the sampled data

## Neural Networks (NN)

- *Analogy*: brain neural connections and learning
- Input layer,  $0 \leq$  hidden layers, output layer
- Domain knowledge helps establish structure
- Model *any* function with enough nodes/layers
- For each layer  $i$  w/  $n_i$  inputs  $\mathbf{a}_i \in \mathbb{R}^{n_i}$ 
  - Bias weight  $b_i \in \mathbb{R}$
  - Input weights to next layer:  $W_i \in \mathbb{R}^{n_i \times n_{i+1}}$
  - Activation function  $g$ 
    - Apply to weighted sum of inputs + bias
    - $\mathbf{a}_{i+1} = g(W_i^T \mathbf{a}_i + b_i)$
- Perceptron (step) activation function
  - $g(z) = 1$  if  $z > 0, 0$  otherwise
- Logistic/Sigmoid Function (Gaussian integrated):
  - $g(z) = \frac{1}{1+e^{-z}}$
- **Back Propagation** learning
  - Initiate input layer weights and bias
  - Observed error:  $L_2(W) = |\mathbf{y} - h_W(\mathbf{x})|^2$
  - Update output weights via observed error

- Use *gradient descent*
- From output to earliest hidden layer:
  - Propagate  $\Delta$  values to previous layer
  - Update weights between the two layers
- (each hidden node resp. for some error)

## Deep Learning

- Applications:
  - B/W image colorization
  - Sound generation from silent media
  - Machine translations
  - Object detection
  - Handwriting generation and recognition
  - [Eloquent] Text generation via recurrent NN
  - Game Playing (in a human-like way)
  - Speech recognition
  - Image object detection, classification, de-abstractization, sketch inversion
  - Natural Language Processing

## Support Vector Machines (SVM)

- Maximum margin separator between examples
- Non-param., but needs frac. of ex. (support vectors)
- Expensive: use quadratic programming -  $O(n^3)$
- Support vectors lie along decision boundary
- **Kernels** (kernel function)
  - notion of similarity between data
  - linearly-separable data:
    - linear kernel:  $K(x, y) = x^T y$
  - non-linearly separable data
    - polynomial:  $K(x, y) = (x^T y + c)^d$
    - project data to higher dimension
    - dividing hyperplane in  $d + 1$  dimensions
      - $d$  degrees of freedom
  - Radial-Basis:  $K(x, y) = e^{-(\|x-y\|^2/(2\sigma^2))}$
  - Sigmoid:  $K(x, y) = \tanh(\alpha x^T y + \theta)$
  - Kernels for other abstract similarity...
    - Requires some domain knowledge of data
- Regularization constant  $C$  and Soft-Margin SVM
  - $\downarrow C$ : permit points to fall inside margin
- Multi-class classification: 1vs1 or 1vs $\forall$  strategies

## Randomized Algorithms/Optimization

- Given i/p  $X$ , objective/fitness func.  $f: X \mapsto \mathbb{R}$ 
  - find  $x \in X$  such that  $f(x) = \max_x f(x)$
- **Hill Climbing (HC)**
  - *Analogy*: Climb Everest in fog w/ amnesia
  - Start at arbitrary location in input space
  - Move to the highest-value neighbor
  - If neighbor  $>$  current, proceed to neighbor
  - Else return current. Possibly only *local max*.
- **Random Restart Hill Climbing (RHC)**
  - Alleviates local max somewhat
  - Restarts at random point a constant # of times.
- **Simulated Annealing (SA, Metropolis Hastings)**
  - *Analogy*: Rept'd heat/cool'g strengthens blade.
  - Allow bad moves, but with decreasing freq.

- $T$ : some gradually decreasing temperature func.
- $\Delta E = \text{Value}(\text{next}) - \text{Value}(\text{curr})$
- If  $\Delta E \leq 0$  take bad move w/ prob.  $e^{\Delta E/T}$
- As  $T \rightarrow 0$ , transitions *Random Walk*  $\rightarrow$  *HC*.

## - Genetic Algorithms (GA)

- *Analogy*: Natural selection and mutation
- *Apps*: Opt problems w/ approp. encoding
  - 8-Queens, circuit layout, job schedule
- Select most fit pairs among population
- Reproduce (cross-over) each pair
  - One-point crossover strategy
  - Uniform crossover strategy (random bits)
- Mutate offspring with small probability
- Replace least-fit individuals with new offspring
- Repeat until convergence

## - MIMIC

- Model probability distribution instead of:
- Population (non-parametric representation)
- Convey structure
- Generate samples from distribution  $P^{\theta_t}(x)$
- Set  $\theta_{t+1}$  to  $n$ 'th percentile
- Retain only samples with  $f(x) \geq \theta_{t+1}$
- Estimate  $\hat{P}^{\theta_{t+1}}(x)$
- Repeat until convergence
- Technical details:
  - Estimate distribution via dependency tree
  - Use the *KL divergence* from info theory
  - Vastly fewer iterations than above algs
  - Each iteration much more costly

## Bayesian

- Most prob.  $h$  given the data:  $\arg\max_{h \in H} \Pr\{h|D\}$
- Bayes Rule:  $\Pr\{h|D\} = \Pr\{D|h\} \Pr\{h\} / \Pr\{D\}$
- Maximum a posteriori (MAP):
  - $h_{MAP} = \arg\max_h \Pr\{D|h\} \Pr\{h\}$
  - disregard the normalizer  $\Pr\{D\}$
- Maximum likelihood:  $h_{ML} = \arg\max_h \Pr\{D|h\}$ 
  - assumes uniform  $\Pr\{h\}$ . Actual  $h$  irrelevant.
  - if Gaussian noise in data  $\rightarrow \arg\min_h L_2$  error.

## - Problems:

- Requires domain knowledge of prior probabilities
- Expensive to compute, being linear in  $|H|$

## - Bayesian Optimal Classifier

- Most probable *classification* given the data
- Combine weighted predictions of  $\forall h_i$  given data
- $V_{MAP} = \arg\max_v \sum_h \Pr\{v|h\} \Pr\{h|D\}$

## - Bayesian Belief Networks (Bayes Nets)

- Cond. indep. assumptions for  $X$  and  $Y$  given  $Z$ 
  - $\Pr\{X|Y, Z\} = \Pr\{X|Z\}$
- Conditional probabilities
- Sample nodes in *Topological Sort*
- *apps*: simulate/approximate a complex process

## - Naive Bayes

- assume data independence given parent.
- $\rightarrow$  assume attr. independence given class.
  - No guarantee in the real world.
- *apps*: classifying text documents
  - Attributes: words. Values: frequencies, OR
  - Attributes: word positions. Values: words.

- $V_{NB} = \arg\max_V \Pr\{V\} \prod_i \Pr\{a_i|V\}$
- Requires sufficiently large set of training data
- Tractable: requires no search

## Classification Metrics

# Unsupervised Algorithms

## Clustering

### - Single-Linkage Clustering

- Start with  $n$  points
- Intercluster distance: closest two points in each
  - Alternatives: median, mean distance
- Merge two closest clusters
- Repeat  $n - k$  times to produce  $k$  clusters
- $O(n^3)$ , but practically fast

### - K-Means Clustering (in Euclidean Space)

- 'Hard' clustering (special case of EM Clustering)
- *Apps*: Image segmentation and compression
- Pick  $k$  centers at random (or distributed)
- Each center "claims" its closest points
  - Closest  $\rightarrow$  minimizing  $L_2$  error.
- Recompute the centers (avg clustered points)
- Repeat until convergence (to *local minimum*!)
- $O(kn)$  per iteration, and  $O(k^n)$  iterations.

### - K-Medoids Clustering

- Not Euclidean-Space  $\rightarrow$  cannot use  $L_2$  error<sub>2</sub>
- Use abstract  $\mathcal{V}(x_n, \mu_k)$  instead of  $\|x_n - \mu_k\|^2$
- Can't take average of clusters.
- Can assign  $\mu_k$  to one of cluster points
- $O(kn + n_k^2)$  per iteration

### - Expectation-Maximization (EM Clustering)

- 'Soft' clustering
  - Point shared by mult. clusters prob'ly
- Gaussian mixture model:
  - $f(x) = \sum_{i=1}^k \pi_i N_i(x|\mu_i, \sigma_i^2)$
  - $\pi_i$ : mixing coefficient.  $\sum_{i=1}^k \pi_i = 1$
- Mixed Gaussian  $(\mu, \sigma^2)$  for each cluster
- Pick initial hidden vars:  $\mu_k$ , covariances, and  $\pi_k$ .
- *E - step*: prob. of component  $k$  explaining  $x$
- *M - step*: Re-estimate hidden vars
- Iterate until log-likelihood reaches convergence
  - $\log \Pr\{\mathbf{X}|\mu, \sigma^2, \pi\} =$ 

$$\sum_{n=1}^N \log \left\{ \sum_{k=1}^K \pi_k N(x_n|\mu_k, \sigma_k^2) \right\}$$
- Local optima possible  $\rightarrow$  random restart.

## Dimensionality Reduction

- *Filtering*: Feature selection abstracted from learner
  - Fast, but ignore the learning problem
  - Ex: DT learning, information gain, variance
- *Wrapping*: Coupled with learner as one routine
  - Takes model bias into account, but SLOW.
  - Ex: Hill Climbing, Rand. Algorithms
- *Forward Search*:
  - Start with empty set of features.
  - Add most 'contributing' features until threshold.

## - Backward Search:

- Start with full set of features.
- Remove most 'useless' features until threshold.

## - Principle Component Analysis (PCA)

- Mutually orthogonal and ordered transformed features
- Gravitate towards 'average' of features
- 'Best' dimension: highest *variance/eigenvalue*
- This is the *First Principle Component*
  - Find via *SVD* in Linear Algebra
- Find orthogonal (*Second Principle*) component
- Idea: min.  $L_2$  error moving from  $N \rightarrow M$  dims
- Select best components with non increasing var.
- Eliminate dimensions beyond best  $M$

## - Independent Component Analysis (ICA)

- Mutually independent transformed features
- Gravitate towards 'abstracted' components
- *Apps*: Cocktail party problem, mixed models
- Min. *mutual information* for transf. features
  - Maximizing *kurtosis* of a dim. is one way.
- $\rightarrow I(y_i; y_j) = 0$  (or minimum)
- $I(X; Y)$  is max. (between *orig.* & transformed)
- Idea: given observables, find indep. hidden var.

## - Random Component Analysis (RCA)

- Linear transformation in random directions
- *Random* linear. combo of orig. dims still useful
- For  $N \rightarrow M$ ,  $M$  normally  $\uparrow$  than PCA/ICA.
- Fast. But requires many random trials.

## - Linear Discriminant Analysis (LDA)

- Finds projection that *discriminates* on label.

## Classification Metrics

- Accuracy
- Precision
- Recall
- $F_b$

# Reinforcement Learning

## Markov Decision Processes (MDP)

- *Markov*: Only present matters. Stationary rules.
- $a \in \text{Actions}$ ,  $s, s' \in \text{States}$
- *Model*:  $T(s, a, s') \approx \Pr\{s'|s, a\}$
- *Reward*:  $R(s), R(s, a), R(s, a, s')$  (typically 1st form)
- *Policy*:  $\pi(s) \rightarrow a$ .
  - $\pi^*$ : optimal policy w/ max reward.
- *Infinite horizon*: navigate world forever,  $\infty$  rewards.
- *Finite horizon*: discount  $\gamma$  to incentivise finish.
- $\pi^*(s) = \arg\max_a \sum_{s'} T(s, a, s') U(s')$
- **Bellman Equation**:
  - $U(s) = R(s) + \gamma \max_a \sum_{s'} T(s, a, s') U(s')$
  - $U(s)$ : state utility,  $R(s)$ : state reward
  - $n$  states  $\rightarrow n$  equations,  $n$ , unknowns
  - non-linear due to *max* operation
- **Value Iteration (VI)**
  - Start w/ arbitrary utilities
  - Update based on neighbors

- $\hat{U}_{t+1}(s) = R(s) + \gamma \max_a \sum_{s'} T(s, a, s') \hat{U}_t(s')$
- Repeat until convergence
- Then solve for  $\pi^*$ , straightforward
- Lends itself for parallel computation
- **Policy Iteration (PI)**
  - Start with arbitrary  $\pi_0$
  - Given  $\pi_t$ , calculate  $U_t = U^{\pi_t}$  (follow policy)
    - $U_t(s) = R(s) + \gamma \sum_{s'} T(s, \pi_t(s), s') U_t(s')$
    - The action is fixed from the policy  $\pi_t(s)$
    - $n$  linear eq,  $n$  unknowns.  $\rightarrow$  Linear Algebra.
  - Improve:  $\pi_{t+1} = \operatorname{argmax}_a \sum_{s'} T(s, a, s') U_t(s')$

## Reinforcement Learning (model-free)

- No model (transition probabilities) or rewards
- Given transitions  $\langle s, a, r, s' \rangle$ , learn policy
- **Q-Learning**
  - $Q(s, a) = R(s) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q(s', a')$
  - Derive  $U$  and  $\pi$  from  $Q$ :
    - $U(s) = \max_a Q(s, a)$
    - $\pi(s) = \operatorname{argmax}_a Q(s, a)$
  - Estimate from transitions (w/out  $R$  &  $T$ ):
    - $Q(s, a) \xleftarrow{\alpha_t} r + \gamma \max_{a'} Q(s', a')$
    - $\alpha_t$ : learning rate
    - $\max_{a'} Q(s', a')$ : util. of next state
    - Notation:  $v \leftarrow x = v \leftarrow (1 - \alpha)v + \alpha x$
  - $\hat{Q}$  starts anywhere (results will vary)
  - Following update rule above,  $Q(s, a) \rightarrow Q(s, a)$
  - Visiting  $s, a \propto$  times,  $s' \approx T(s, a, s'), r \approx R(s)$
  - Questions:
    - Initial  $\hat{Q}$
    - $\alpha_t$  decay factor?
    - action choosing policy?
      - Take random action *sometimes*
      - Otherwise take best action
      - Mimics simulated annealing
  - *Exploration vs. Exploitation*
    - Exploration: cont. to learn  $Q$
    - Exploitation: quicker maximize  $\pi^*$

- A uses strategy 1 w/ probability  $p$
- $eg_i$ : Expected gain for A given B uses strategy  $i$
- $eg_1 = m_{11}p + m_{21}(1 - p)$
- $eg_2 = m_{12}p + m_{22}(1 - p)$
- $p^* = \max_p \{ \min \{ eg_1, eg_2 \} \}$ 
  - Either  $p = 0, p = 1$ , or where both equal
- Optimal *Mixed Strategy* for  $n \times m$  game
  - $\mathbf{P}_A = \{p_1, p_2, \dots, p_n\}$
  - $eg_j = \sum_{i=1}^n m_{ij} p_i$  for all  $j \in 1, \dots, m$
  - $n$  unknowns,  $m$  eq.  $\rightarrow$  Linear Programming
  - Choose  $\mathbf{P}_A$  to maximize  $\min \{ eg_1, \dots, eg_m \}$
  - such that  $\sum p_i = 1, 0 \leq p_i \leq 1, \forall i$

## Non-Zero-Sum Games

- Non-deterministic, non-cooperative, hidden info
- Typical example: Prisoner's Dilemma
- $n$  strategy spaces  $\mathbf{S}_1, \mathbf{S}_2, \dots, \mathbf{S}_n$
- $n$  payoff functions  $u_1, \dots, u_n$  st  $u_i : \mathbf{S}_1 \times \dots \times \mathbf{S}_n \rightarrow \mathbb{R}$
- What's the optimal mixed strategy?
- **Nash Equilibrium (NE)**
  - $s_1^*, \dots, s_n^* \in S_1 \times \dots \times S_n$  are a NE iff
    - $\forall i. s_i^* = \operatorname{argmax}_{s_i} u_i(s_1^*, \dots, s_i, \dots, s_n^*)$
    - ( $\neg \exists$  reason for any one player to switch)
  - Works for pure & mixed strategies
  - Assumes best action for self (regardless of others)
  - Can contain multiple Nash Equilibria
  - An *implausible* threat hinders own utility
  - $n$  repeated games  $\rightarrow n$  repeated NE
    - assuming implausible threats

## Bayesian Games

- Action spaces:  $\mathbf{A}_1, \dots, \mathbf{A}_n$
- Type spaces:  $\mathbf{T}_1, \dots, \mathbf{T}_n$
- Beliefs:  $\mathbf{P}_1, \dots, \mathbf{P}_n$
- $P_{-i}(t_{-i}|t_i) = \text{PDF of others' types given own}$
- Payoff functions:  $u_1, \dots, u_n$
- $u_i(a_1, \dots, a_n, t_i)$  - payout to player  $i$  w/ type  $i$ 
  - player  $j$  chooses action  $a_j$  for  $\forall j$
- All players know their own  $\mathbf{A}_i, \mathbf{T}_i, \mathbf{P}_i, u_i$
- Strategy  $S_i : \mathbf{T}_i \rightarrow \mathbf{A}_i$
- **Bayesian Nash Equilibrium (BNE)**
  - $s_1^*, \dots, s_n^* \in S_1 \times \dots \times S_n$  are a BNE iff  $\forall i$  and  $\forall t_i \in \mathbf{T}_i. s_i^*(t_i) = \operatorname{argmax}_{a_i \in \mathbf{A}_i} \left\{ \sum_{t_{-i} \in \mathbf{T}_{-i}} u_i(s_1^*(t_1), \dots, a_i, \dots, s_n^*(t_n)) \times P_i(t_{-i}|t_i) \right\}$

## Repeated games w/ uncertain end

- # rounds left is uncertain
- $\Pr \{ \text{play again} \} = \gamma, \Pr \{ \text{game over} \} = 1 - \gamma$
- $E[\# \text{ rounds}] = 1/(1 - \gamma)$
- **Tit for Tat** strategy

- Cooperate 1st round, copy opponent move after
- **Grim Trigger** strategy
  - Cooperate while opponent cooperates.
  - Once line is crossed, forever defect
- **Pavlov** strategy
  - Cooperate if agree, defect if disagree.
- *Subgame perfect (SP)*
  - Str. is always best response indep. of history
  - $\neg \text{SP}$  if  $\exists$  history of moves st. str. implausible
  - $\neg \text{SP} \iff \exists$  pair of str's not leading back to mutual cooperate
- *Mini-Max Profile* (For zero-sum game)
  - Min guaranteed payoffs for  $\forall$  players on defense
  - Applicable in pure or mixed strategies

## Stochastic Games & Multiagent RL

- *Analogy*: MDP-RL::Stochastic game-Multiagent RL
- More general than MDPs or other previous models
- **S**: states
- **A<sub>i</sub>**: actions for player  $i, a, b, a \in \mathbf{A}_1, b \in \mathbf{A}_2$
- **T**: transition probabilities  $T(s, (a, b), s')$
- **R<sub>i</sub>**: rewards for payer  $i, R_1(s, (a, b)), R_2(s, (a, b))$
- $\gamma$ : discount
- Impose restrictions to produce other models:
  - $R_1 = -R_2 \rightarrow$  0-sum stochastic game
  - $T(s, (a, b), s') = T(s, (a, b'), s'), R_2(s, (a, b)) = 0, R_1(s, (a, b)) = R_1(s, (a, b'))$
  - $\forall b' \rightarrow$  MDP, makes second player irrelevant
  - $|\mathbf{S}| = 1 \rightarrow$  repeated games
- **Zero-Sum Stochastic Games**
  - $Q_i^*(s, (a, b)) = R_i(s, (a, b)) + \gamma \sum_{s'} T(s, (a, b), s') \min_{a', b'} Q_i^*(s', (a', b'))$
  - *Q-Learning* for transition  $\langle s, (a, b), (r_1, r_2), s' \rangle$ 
    - $Q_i(s, (a, b)) \xleftarrow{\alpha} r_i + \gamma \min_{a', b'} Q_i(s', (a', b'))$
    - Value Iteration works
    - Minimax-Q converges
    - Unique solution to  $Q^*$
    - Policies can be computed independently
    - Update efficient (polytime)
    - Q functions sufficient to specify policy
- **General-Sum Stochastic Games (Nash-Q)**
  - Same as  $Q_i^*/Q$ -Update above, but use Nash Equilibrium instead of minimax/maximin
  - Value Iteration doesn't work
  - Nash-Q doesn't work
  - No unique solution
  - Policies can't be computed independently
  - Update not efficient
  - Q functions not sufficient to specify policy

## Game Theory

### Zero-Sum Games

- **Minimax/Maximin**
  - Player A considers worse-case strategy by B
  - A chooses maximum minimum value by B
  - B chooses the minimum maximum value by A
  - Applies to *perfect* or *hidden* information games
- *Von Neumann*
  - 0-sum games of per. info: *minimax* = *maximin*.
  - $\exists$  optimal strategy for each player
- Optimal *Mixed Strategy* for 2X2 game
  - Two players, A & B, strategies  $A_1, A_2, B_1, B_2$
  - Mixed strategy gains:  $\mathbf{m}_{AB} : m_{11}, m_{21}, m_{12}, m_{22}$